

Investigate_a_Dataset

October 7, 2020

1 Project 2: Investigating Data From Around the World

1.1 Introduction - Gapminder Data

In the below analysis I'll be driving into some topics that sparked my interest. The data from Gapminder was very extensive with lots of datasets to be explored but I was able to narrow it down to a few topics. I would like to preface, since each csv file downloaded from Gapminder had varying levels of completion and available years showing data I can't combine all my data into one dataset for all questions. Instead, I'll be analyzing, wrangling and cleaning the data to create specific dataframes for multiple research questions. Each question will be broken into sections: question, explanation/information, cleaning/wrangling, analysis and synopsis.

1.2 Project Introduction

Throughout this project our questioning will revolve around 2 distinct variables: total murders per year and life expectancy. Part 1 (Q1 & Q2) help us understand the correlation between total murders and population. We will examine this topic further by representing the most dangerous countries by violent crime rates as a proportion to population size and population density. Part 2 (Q3 & Q4) we will understand how average life expectancy changes throughout history and if income per capita or health spend based on GDP has an effect towards more prolonged life. Finally, I will end the project with a final conclusion summarizing our findings and data limitations.

Here are the 4 research questions we are about to explore: 1. Where are the most dangerous places to live? 2. How does population density effect violent murder rates? 3. Does average income per person effect life expectancy by country and which country had the greatest improvement in life expectancy? 4. How health spend effect life expectancy?

1.3 Table of Contents

Datasets Overview

Largest Challenge

Exploratory Data Analysis

Q1: Dangerous Living

Q2: Murders vs Density

Q3: Most Improved & Life Expectancy among Wealth

Q4: Health Spends vs Life Expectancy

Final Conclusion

In [1]: *#import packages - I have used these imports at the start of most questions as well just*

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
% matplotlib inline
```

Dataset Overview **NOTE:** All of these datasets are columnized by year and organized vertically by country. The data points correspond to a country for a one year period. Also, all datasets have variety of years and provided countries.

What the data values represent: **income.csv** - Income per person (GDP per capita, PPP inflation-adjusted) **life_exp.csv** - Average life expectancy in years **murders.csv** - Total murders **total_health_spend.csv** - Total health spend (% of total GDP) **pop_density.csv** - Average number of people per square km of land in country **populations.csv** - Total population **gini.csv** - Gini coefficient per year (Higher number greater income inequality). Used for example

In [2]: *#How I read in all datasets throughout the project*

```
gini_df = pd.read_csv('gini.csv')
income_df = pd.read_csv('income.csv')
life_exp_df = pd.read_csv('life_exp.csv')
murders_df = pd.read_csv('murders.csv')
pop_density_df = pd.read_csv('pop_density.csv')
pop_df = pd.read_csv('population.csv')
health_df = pd.read_csv('total_health_spend.csv')
```

Largest Challenge in Project

All the datasets from Gapminder organized their data using years as columns for each corresponding country and listed irrelevant future years in some dataframes. In order to analyze the data properly over time I needed to consolidate all the year columns into one 'year' column and eliminate the irrelevant years. I did this by using drop function for irrelevant years and the melt function to organize my data as such: country, year, data1, data2, etc.

Below are the corresponding steps needed to complete the melt and drop function. A method I used multiple times throughout the project.

In [3]: *#Understand data structure - gives me what years I need to drop*

```
gini_df.head()
```

```
Out[3]:
```

	country	1800	1801	1802	1803	1804	1805	1806	1807	1808	...	\
0	Afghanistan	30.5	30.5	30.5	30.5	30.5	30.5	30.5	30.5	30.5	...	
1	Albania	38.9	38.9	38.9	38.9	38.9	38.9	38.9	38.9	38.9	...	

2	Algeria	56.2	56.2	56.2	56.2	56.2	56.2	56.2	56.2	56.2	...
3	Andorra	40.0	40.0	40.0	40.0	40.0	40.0	40.0	40.0	40.0	...
4	Angola	57.2	57.2	57.2	57.2	57.2	57.2	57.2	57.2	57.2	...

	2031	2032	2033	2034	2035	2036	2037	2038	2039	2040
0	36.8	36.8	36.8	36.8	36.8	36.8	36.8	36.8	36.8	36.8
1	29.0	29.0	29.0	29.0	29.0	29.0	29.0	29.0	29.0	29.0
2	27.6	27.6	27.6	27.6	27.6	27.6	27.6	27.6	27.6	27.6
3	40.0	40.0	40.0	40.0	40.0	40.0	40.0	40.0	40.0	40.0
4	42.6	42.6	42.6	42.6	42.6	42.6	42.6	42.6	42.6	42.6

[5 rows x 242 columns]

In [4]: gini_df.shape

Out[4]: (195, 242)

```
In [5]: #Get rid of years above 2020
col = range(2021, 2041)
for i in col:
    col_drop = str(i)
    gini_df.drop(col_drop, axis=1, inplace=True)
```

```
In [6]: #Makes sure code ran properly and dataset is updated
gini_df.head()
```

Out[6]:

	country	1800	1801	1802	1803	1804	1805	1806	1807	1808	...	\
0	Afghanistan	30.5	30.5	30.5	30.5	30.5	30.5	30.5	30.5	30.5	...	
1	Albania	38.9	38.9	38.9	38.9	38.9	38.9	38.9	38.9	38.9	...	
2	Algeria	56.2	56.2	56.2	56.2	56.2	56.2	56.2	56.2	56.2	...	
3	Andorra	40.0	40.0	40.0	40.0	40.0	40.0	40.0	40.0	40.0	...	
4	Angola	57.2	57.2	57.2	57.2	57.2	57.2	57.2	57.2	57.2	...	

	2011	2012	2013	2014	2015	2016	2017	2018	2019	2020
0	36.8	36.8	36.8	36.8	36.8	36.8	36.8	36.8	36.8	36.8
1	29.3	29.1	29.0	29.0	29.0	29.0	29.0	29.0	29.0	29.0
2	28.2	27.9	27.7	27.6	27.6	27.6	27.6	27.6	27.6	27.6
3	40.0	40.0	40.0	40.0	40.0	40.0	40.0	40.0	40.0	40.0
4	42.7	42.6	42.6	42.6	42.6	42.6	42.6	42.6	42.6	42.6

[5 rows x 222 columns]

For more information on [pd.melt](#)

```
In [7]: #Reorganize my dataset to make it more easily analyzed columns: (country, year, gini)
gini_df2 = pd.melt(gini_df, id_vars='country', var_name='year', value_name='gini')

#Check function ran properly
gini_df2.head()
```

```
Out[7]:
```

	country	year	gini
0	Afghanistan	1800	30.5
1	Albania	1800	38.9
2	Algeria	1800	56.2
3	Andorra	1800	40.0
4	Angola	1800	57.2

```
In [8]: gini_df2.shape
```

```
Out[8]: (43095, 3)
```

```
In [9]: #Understand dataset
income_df.head()
```

```
Out[9]:
```

	country	1800	1801	1802	1803	1804	1805	1806	1807	1808	...	\
0	Afghanistan	603	603	603	603	603	603	603	603	603	...	
1	Albania	667	667	667	667	667	668	668	668	668	...	
2	Algeria	715	716	717	718	719	720	721	722	723	...	
3	Andorra	1200	1200	1200	1200	1210	1210	1210	1210	1220	...	
4	Angola	618	620	623	626	628	631	634	637	640	...	

	2031	2032	2033	2034	2035	2036	2037	2038	2039	2040
0	2550	2600	2660	2710	2770	2820	2880	2940	3000	3060
1	19400	19800	20200	20600	21000	21500	21900	22300	22800	23300
2	14300	14600	14900	15200	15500	15800	16100	16500	16800	17100
3	73600	75100	76700	78300	79900	81500	83100	84800	86500	88300
4	6110	6230	6350	6480	6610	6750	6880	7020	7170	7310

[5 rows x 242 columns]

```
In [10]: #Drop years after 2020
col1 = range(2021, 2041)
for i in col1:
    col_drop = str(i)
    income_df.drop(col_drop, axis=1, inplace=True)

#Melt the columns to rows
income_df_updated = pd.melt(income_df, id_vars='country', var_name='year', value_name=

In [11]: #Check updates
income_df_updated.head()
```

```
Out[11]:
```

	country	year	income
0	Afghanistan	1800	603
1	Albania	1800	667
2	Algeria	1800	715
3	Andorra	1800	1200
4	Angola	1800	618

Research Question 1: Where are the most dangerous places to live?

```

In [12]: #Import packages:
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
% matplotlib inline

In [13]: #Datasets needed:
murders_df = pd.read_csv('murders.csv')
pop_df = pd.read_csv('population.csv')

```

1.3.1 Explanation

I'm using total murders as the metric to decide how dangerous a country is. If there is a higher murder count in a give country than we can assume it is more dangerous. Of course, we can argue that if the population is greater than total murders may be higher in comparison to smaller countries. Therefore, I'll create an additional column using murders in proportion to the countries population.

1.3.2 Understanding Data Layout

```

In [14]: #Understand dataframe layout - we immediately see some irrelevant future years
murders_df.head()

```

```

Out[14]:
   country  1990  1991  1992  1993  1994  1995  1996 \
0  Afghanistan  2070.00  2200.00  2380.00  2600.00  2830.00  3020.00  3160.00
1    Albania    160.00   182.00   201.00   221.00   239.00   267.00   295.00
2    Algeria    377.00   382.00   391.00   400.00   410.00   421.00   432.00
3   Andorra     0.48    0.51    0.54    0.55    0.55    0.53    0.51
4    Angola    527.00   532.00   543.00   569.00   598.00   608.00   582.00

   1997  1998  ...  2007  2008  2009  2010  2011 \
0  3270.0  3350.00  ...  4910.00  4960.00  4990.00  4940.00  5020.00
1   327.0   338.00  ...    82.70   77.50   67.50   68.40   68.50
2   452.0   468.00  ...   437.00  441.00  445.00  447.00  451.00
3    0.5    0.49  ...    0.52   0.52   0.53   0.54   0.54
4   582.0   667.00  ...   904.00  933.00  958.00  978.00  990.00

   2012  2013  2014  2015  2016
0  5190.00  5560.00  5820.00  6060.00  6270.00
1    68.50   68.70   68.90   69.20   69.50
2   457.00   465.00   474.00   472.00   471.00
3    0.54    0.55    0.55    0.55    0.55
4  1010.00  1030.00  1050.00  1080.00  1090.00

[5 rows x 28 columns]

```

```

In [15]: #Compare shapes to understand which dataframe we have to match to the other
murders_df.shape

```

```
Out[15]: (187, 28)
```

```
In [16]: pop_df.shape
```

```
Out[16]: (195, 302)
```

We notice, murders dataset has a smaller range of years (27 vs 301) and listed countries (187 vs 195) than population. **### Cleaning/Wrangling Data**

Steps: 1. Remove irrelevant years from population to match murders dataframe 2. Melt both dataframes for easier analysis 3. Merge dataframes only keeping similar countries 4. Check for missing or null values

```
In [17]: #Drop years - need 1990 - 2016
col = np.r_[1800:1990 , 2017:2101]
for i in col:
    col_drop = str(i)
    pop_df.drop(col_drop, axis=1, inplace=True)

#Check for code completion
pop_df.head()
```

```
Out[17]:
```

	country	1990	1991	1992	1993	1994	1995	\
0	Afghanistan	12400000	13300000	14500000	15800000	17100000	18100000	
1	Albania	3290000	3280000	3250000	3200000	3150000	3110000	
2	Algeria	25800000	26400000	27000000	27600000	28200000	28800000	
3	Andorra	54500	56700	58900	61000	62700	63900	
4	Angola	11800000	12200000	12700000	13100000	13500000	13900000	

	1996	1997	1998	...	2007	2008	2009	\
0	18900000	19400000	19700000	...	27100000	27700000	28400000	
1	3100000	3100000	3110000	...	3030000	3000000	2970000	
2	29300000	29700000	30200000	...	34200000	34700000	35300000	
3	64400	64300	64100	...	82700	83900	84500	
4	14400000	14900000	15400000	...	20900000	21700000	22500000	

	2010	2011	2012	2013	2014	2015	2016	
0	29200000	30100000	31200000	32300000	33400000	34400000	35400000	
1	2950000	2930000	2910000	2900000	2900000	2890000	2890000	
2	36000000	36700000	37400000	38100000	38900000	39700000	40600000	
3	84500	83700	82400	80800	79200	78000	77300	
4	23400000	24200000	25100000	26000000	26900000	27900000	28800000	

[5 rows x 28 columns]

```
In [18]: #Melt both datasets so we can join
murders_melt_df = pd.melt(murders_df, id_vars='country', var_name='year', value_name='t
pop_melt_df = pd.melt(pop_df, id_vars='country', var_name='year', value_name='total_pop
```

```

In [19]: #Check to make sure there are 3 columns
murders_melt_df.shape

Out[19]: (5049, 3)

In [20]: pop_melt_df.shape

Out[20]: (5265, 3)

In [21]: #Merge datasets and check for null values
violence_df = pd.merge(pop_melt_df, murders_melt_df, on=['country', 'year'], how='right')
violence_df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 5049 entries, 0 to 5048
Data columns (total 4 columns):
country          5049 non-null object
year             5049 non-null object
total_pop        5049 non-null int64
total_murders    5049 non-null float64
dtypes: float64(1), int64(1), object(2)
memory usage: 197.2+ KB

In [22]: violence_df.head()

Out[22]:
```

	country	year	total_pop	total_murders
0	Afghanistan	1990	12400000	2070.00
1	Albania	1990	3290000	160.00
2	Algeria	1990	25800000	377.00
3	Andorra	1990	54500	0.48
4	Angola	1990	11800000	527.00

According to the above info and head functions we can see that there are no null values and the dataframes are combined appropriately.

1.3.3 Analysis

In order to compare violent murders and population throughout time I need to group our data points by country and find the mean murders value. This will organize the data allowing me to simply compare total murders based on country. **Note:** Normally, I might create different sections of time to see how murders and population has changed overtime in comparison to countries. Since we are only using a 26 year time span here I decided not to. We will see sectioned time eras later in project

```

In [23]: #This is a global function used to eliminate scientific notation. Also allows multiple
pd.set_option('display.float_format', lambda x: '%.5f' %x)

#Group based on country and find mean from 1990 - 2016
violence_mean_df = violence_df.groupby('country', as_index=False).mean().sort_values(by=
#Top 10 countries with highest
violence_mean_df.head(10)

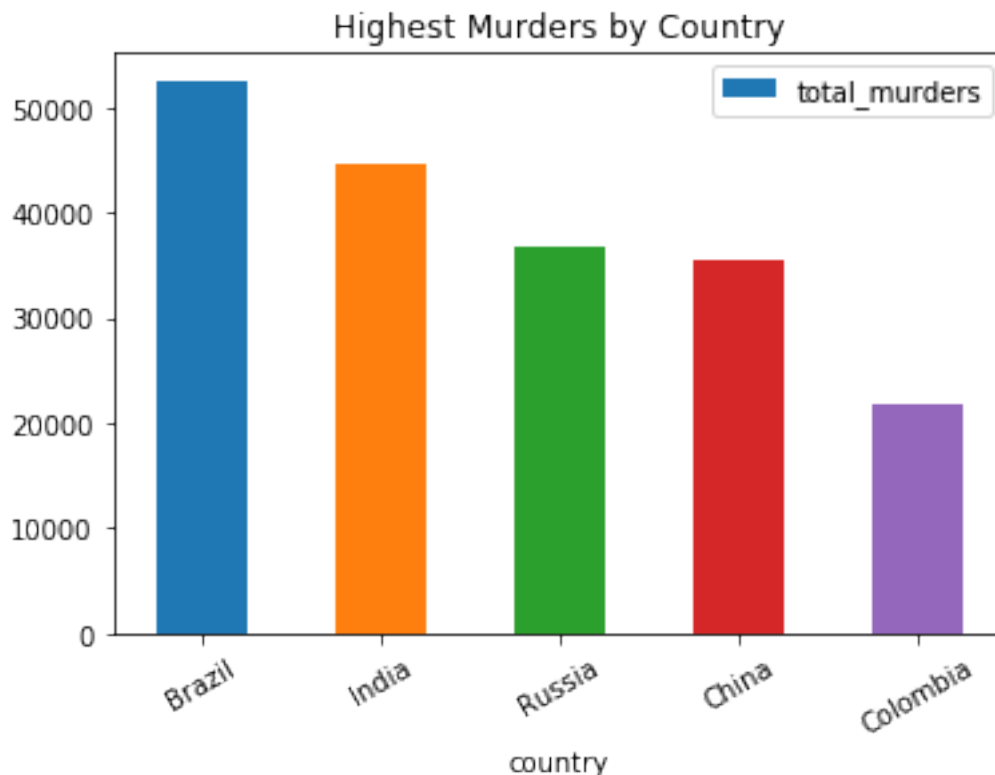
```

```
Out[23]:
```

	country	total_pop	total_murders
23	Brazil	180074074.07407	52555.55556
75	India	1107481481.48148	44740.74074
137	Russia	145555555.55556	36651.85185
35	China	1308888888.88889	35429.62963
36	Colombia	41111111.11111	21848.14815
151	South Africa	46681481.48148	20481.48148
178	United States	288814814.81481	19081.48148
107	Mexico	103488888.88889	16796.29630
132	Philippines	82848148.14815	11578.88889
182	Venezuela	25370370.37037	8288.88889

```
In [24]: #Quick visual representation of countries with highest total murders
violence_mean_df.head().plot(x='country', y='total_murders', kind='bar', title='Highest
```

```
Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x7f64d24613c8>
```



As mentioned before, population of the above countries is much higher than other countries. We can argue that of course there will be more murders merely because there is a greater population, all other things held constant.

How to Fix: We create a column in table using murders in proportion to countries population (total murders / total population)


```
In [25]: #Let's now take into consideration popluation
violence_mean_df['murders_per_capita'] = violence_mean_df['total_murders'] / violence_m


#Check for proper code completion and sort to see top 10 coutries with highest murder r
violence_mean_df.sort_values(by='murders_per_capita', ascending=False).head(10)
```

```
Out[25]:
```

	country	total_pop	total_murders	murders_per_capita
52	El Salvador	5925925.92593	3276.29630	0.00055
36	Colombia	41111111.11111	21848.14815	0.00053
151	South Africa	46681481.48148	20481.48148	0.00044
67	Guatemala	12658888.88889	4494.44444	0.00036
182	Venezuela	25370370.37037	8288.88889	0.00033
93	Lesotho	1964074.07407	595.92593	0.00030
23	Brazil	180074074.07407	52555.55556	0.00029
72	Honduras	7101481.48148	2064.81481	0.00029
137	Russia	145555555.55556	36651.85185	0.00025
160	Swaziland	1002222.22222	208.66667	0.00021

```
In [26]: #Cleaner visual representation of data answering our initial question
x = violence_mean_df.sort_values(by='murders_per_capita', ascending=False)['country'].h
y = violence_mean_df.sort_values(by='murders_per_capita', ascending=False)['murders_per
plt.bar(x, y, color=['Crimson', 'Red', 'Orangered', 'Orange', 'Gold'])
plt.xlabel('Country', fontsize=14)
plt.ylabel('Murders per Capita', fontsize=14)
plt.title('Highest Murders per Capita by Country', fontsize=18)
```

```
Out[26]: Text(0.5,1,'Highest Murders per Capita by Country')
```



1.3.4 Q1 Synopsis

Based on the murders per capita and only taking into consideration the countries that we had data for in the above dataframes provided by Gapminder, Colombia, El Salvador, Guatemala, South Africa and Venezuela are the most dangerous places to live. Of course, this is only based on a few parameters and I'm sure there are many other factors that play into which country is most dangerous.

Possibly population density rather than total population?

Research Question 2: How does population density effect violent murder rates?

This is a continuation of above question so we're able to use the already cleaned and wrangled data. All we need to do is clean/wrangle and then combine population density to total murders.

1.3.5 Explanation

Density is average number of people in square km of land. The reason I'm interested in population density is what if a smaller country with a very large density has a correlation to total murders in the country. We can show correlation visually by using a scatter plot. If total murders increase as population density increases then there is a correlation between the variables.

1.3.6 Clean/Wrangling

```
In [27]: # Continuation of above research question. Now lets add population density to the dataf
pop_density_df = pd.read_csv('pop_density.csv')
pop_density_df.head()
```

```
Out[27]:
```

	country	1950	1951	1952	1953	1954	1955	1956	\
0	Afghanistan	11.90000	12.00000	12.20000	12.30000	12.50000	12.70000	12.90000	
1	Albania	46.10000	47.00000	48.00000	49.20000	50.50000	51.80000	53.30000	
2	Algeria	3.73000	3.79000	3.86000	3.93000	4.01000	4.10000	4.20000	
3	Andorra	13.20000	14.20000	15.40000	16.70000	18.20000	19.60000	21.20000	
4	Angola	3.65000	3.70000	3.78000	3.87000	3.96000	4.05000	4.12000	
		1957	1958	...	2091	2092	2093	2094	\
0		13.10000	13.30000	...	117.00000	117.00000	117.00000	117.00000	
1		54.80000	56.30000	...	48.50000	47.50000	46.50000	45.50000	
2		4.31000	4.41000	...	29.60000	29.60000	29.60000	29.60000	
3		22.90000	24.70000	...	134.00000	134.00000	134.00000	134.00000	
4		4.19000	4.26000	...	135.00000	136.00000	138.00000	140.00000	
		2095	2096	2097	2098	2099	2100		
0		116.00000	116.00000	116.00000	115.00000	115.00000	115.00000		
1		44.50000	43.50000	42.50000	41.60000	40.60000	39.70000		
2		29.70000	29.70000	29.70000	29.70000	29.70000	29.70000		
3		133.00000	133.00000	133.00000	133.00000	133.00000	133.00000		
4		142.00000	144.00000	146.00000	147.00000	149.00000	151.00000		

[5 rows x 152 columns]

```
In [28]: #Delete none used years (1950 - 1989 & 2017 - 2100)
col = np.r_[1950:1990, 2017:2101]
for i in col:
    col_drop = str(i)
    pop_density_df.drop(col_drop, axis=1, inplace=True)

#Melt the columns to rows
pop_density_melt_df = pd.melt(pop_density_df, id_vars='country', var_name='year', value

#Merge into violence_df
violence_density_df = pd.merge(violence_df, pop_density_melt_df, on=['country', 'year'])

#Adding murders_per_capita for fun of comparison
violence_density_df['murders_per_capita'] = violence_density_df['total_murders'] / viol

#Check for null values and shape - should be no null and 5049 items
violence_density_df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 5049 entries, 0 to 5048
Data columns (total 6 columns):
country          5049 non-null object
year             5049 non-null object
total_pop        5049 non-null int64
total_murders    5049 non-null float64
density          5049 non-null float64
murders_per_capita 5049 non-null float64
dtypes: float64(3), int64(1), object(2)
memory usage: 276.1+ KB
```

```
In [29]: #double check data layout
violence_density_df.head()
```

```
Out[29]:
```

	country	year	total_pop	total_murders	density	murders_per_capita
0	Afghanistan	1990	12400000	2070.00000	19.00000	0.00017
1	Albania	1990	3290000	160.00000	120.00000	0.00005
2	Algeria	1990	25800000	377.00000	10.80000	0.00001
3	Andorra	1990	54500	0.48000	116.00000	0.00001
4	Angola	1990	11800000	527.00000	9.50000	0.00004

1.3.7 Analysis

```
In [30]: #Group data by country and find the mean over time. Again, we are only working 26 year
violence_density_mean = violence_density_df.groupby('country', as_index=False).mean()
```

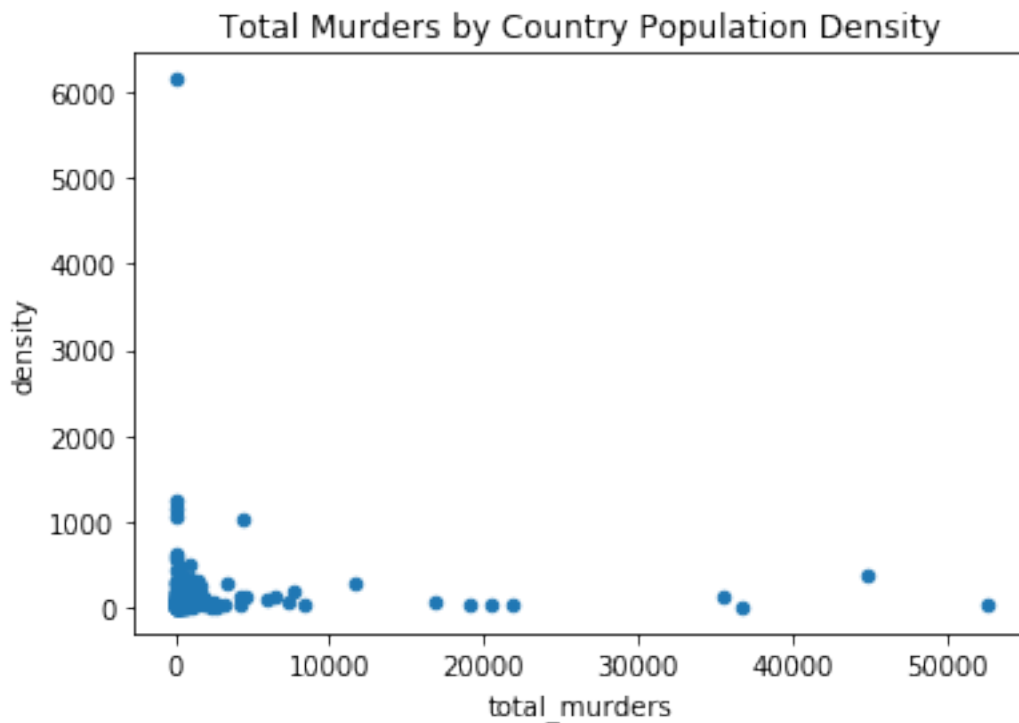
```
#Dataframe sorted by top density
violence_density_mean.sort_values(by='density', ascending=False).head()
```

```
Out[30]:
```

	country	total_pop	total_murders	density	murders_per_capita
146	Singapore	4306296.29630	28.29630	6151.48148	0.00001
103	Malta	398962.96296	5.55333	1246.29630	0.00001
12	Bahrain	880666.66667	17.02593	1157.81481	0.00002
101	Maldives	317629.62963	3.14963	1059.59259	0.00001
13	Bangladesh	132814814.81481	4332.22222	1020.14815	0.00003

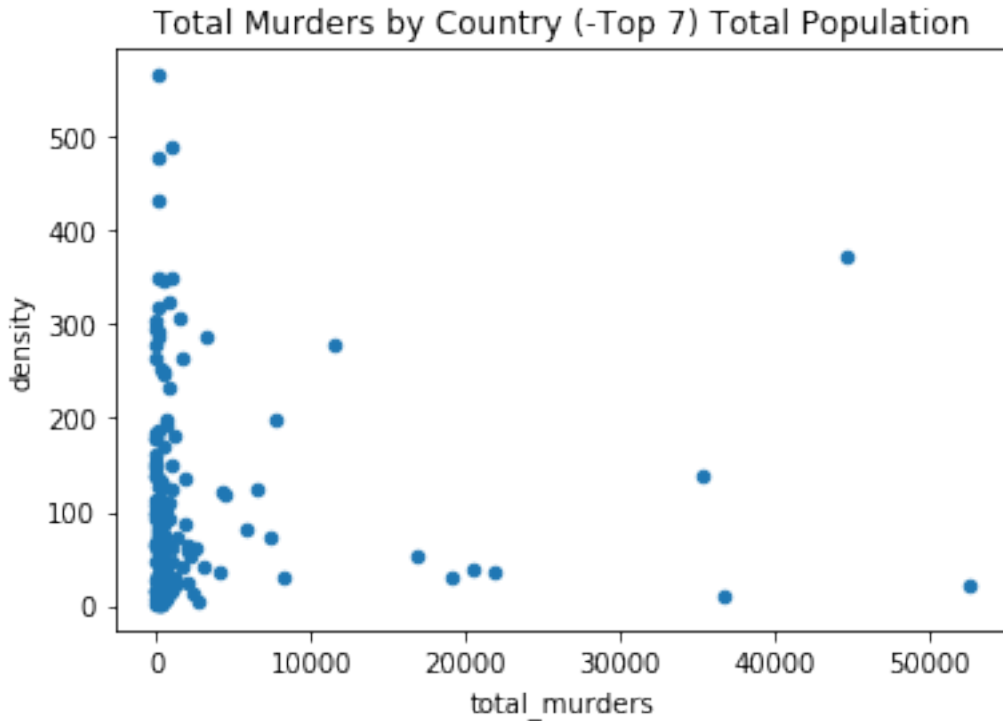
```
In [31]: #Create a scatter plot to show comparison
violence_density_mean.sort_values(by='density', ascending=False).plot(x='total_murders'
```

```
Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x7f64d24815c0>
```



```
In [32]: #I removed the top 7 densest countries because they were outliers and screwing the graph
violence_density_mean.sort_values(by='density', ascending=False).tail(180).plot(x='total_murders'
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x7f64cfb9c630>
```



In order to be absolute certain, I will use the `.corr` function with the Pearson method to find the correlation coefficient between `total_murders` and `density`.

```
In [33]: column_1 = violence_density_mean["total_murders"]
         column_2 = violence_density_mean["density"]
         correlation = column_1.corr(column_2)
         print(correlation * 100)
```

```
-2.54509891184
```

1.3.8 Q2 Synopsis

There is no correlation to murders and population density. Singapore is a great example. Singapore has the highest density for population but has one of the lowest violent murder rates per capita (see below - murder per capita 55X lower than Colombia and density of 166X higher).

Furthermore, there is only a -2.55% correlation between population density and `total_murders` between countries. Using Pearson's method coefficient we would need to see upwards of (-)70% correlation to make an argument for comparison.

Again, the population density is based on the square km area of the entire country. It would probably be a better comparison if we had data for top metropolitan areas in all countries as some countries have a ton of uninhabited area.

```
In [34]: violence_density_mean.query('country == "Singapore" | country == "Colombia"')
```

```
Out [34]:
```

	country	total_pop	total_murders	density	murders_per_capita
36	Colombia	41111111.11111	21848.14815	37.03333	0.00055
146	Singapore	4306296.29630	28.29630	6151.48148	0.00001

Research Question 3: Does average income per person effect life expectancy by country and which country had the greatest improvement in life expectancy?

1.3.9 Information

What is PPP? PPP or Purchasing Power Parity is a way to compare the purchasing power from one country to another. This is also known as international currency and the unit being used in income.csv file. **How do we convert PPP to USD?** Well, we don't. The PPP is based off the purchasing power from other countries to the United States using the USD as a reference. Since we'd be comparing USD to USD the ratio from PPP to USD is 1/1.

1.3.10 Explanation

I was very curious as to how income effects the life expectancy in countries. Below I will explore this more but here is a bit of an overview as the dataframes are a bit more extensive. > As I explored the dataframes I was able to put together a greater range of years to compare. Meaning that I'm not going to be able to take a mean over the entire time period as the data maybe skewed. **To fix this** I created 50 year segments and created another table to better understand life expectancy overtime.

Furthermore, I would like to preface that the years of life expectancy in eariler years was very low and makes me question the integrity of the data in many countries. We are going to assume for the sake of the assignment that the data from Gapminder is percise.

```
In [35]: #Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
% matplotlib inline

In [36]: #Datasets needed:
income_df = pd.read_csv('income.csv')
life_exp_df = pd.read_csv('life_exp.csv')

In [37]: #Understand data layout
income_df.head()
```

```
Out [37]:
```

	country	1800	1801	1802	1803	1804	1805	1806	1807	1808	...	\
0	Afghanistan	603	603	603	603	603	603	603	603	603	...	
1	Albania	667	667	667	667	667	668	668	668	668	...	
2	Algeria	715	716	717	718	719	720	721	722	723	...	
3	Andorra	1200	1200	1200	1200	1210	1210	1210	1210	1220	...	
4	Angola	618	620	623	626	628	631	634	637	640	...	
		2031	2032	2033	2034	2035	2036	2037	2038	2039	2040	

0	2550	2600	2660	2710	2770	2820	2880	2940	3000	3060
1	19400	19800	20200	20600	21000	21500	21900	22300	22800	23300
2	14300	14600	14900	15200	15500	15800	16100	16500	16800	17100
3	73600	75100	76700	78300	79900	81500	83100	84800	86500	88300
4	6110	6230	6350	6480	6610	6750	6880	7020	7170	7310

[5 rows x 242 columns]

In [38]: `life_exp_df.head()`

```
Out[38]:
```

	country	1800	1801	1802	1803	1804	1805	1806	\
0	Afghanistan	28.20000	28.20000	28.20000	28.20000	28.20000	28.20000	28.10000	
1	Albania	35.40000	35.40000	35.40000	35.40000	35.40000	35.40000	35.40000	
2	Algeria	28.80000	28.80000	28.80000	28.80000	28.80000	28.80000	28.80000	
3	Andorra	nan	nan	nan	nan	nan	nan	nan	
4	Angola	27.00000	27.00000	27.00000	27.00000	27.00000	27.00000	27.00000	

	1807	1808	...	2091	2092	2093	2094	2095	\
0	28.10000	28.10000	...	76.50000	76.60000	76.70000	76.90000	77.00000	
1	35.40000	35.40000	...	87.40000	87.50000	87.60000	87.70000	87.80000	
2	28.80000	28.80000	...	88.30000	88.40000	88.50000	88.60000	88.70000	
3	nan	nan	...	nan	nan	nan	nan	nan	
4	27.00000	27.00000	...	78.70000	78.90000	79.00000	79.10000	79.30000	

	2096	2097	2098	2099	2100
0	77.10000	77.30000	77.40000	77.50000	77.70000
1	87.90000	88.00000	88.10000	88.20000	88.30000
2	88.80000	88.90000	89.00000	89.10000	89.20000
3	nan	nan	nan	nan	nan
4	79.40000	79.50000	79.70000	79.80000	79.90000

[5 rows x 302 columns]

In [39]: *#Understand and compare shape*
`income_df.shape`

Out[39]: (193, 242)

In [40]: `life_exp_df.shape`

Out[40]: (187, 302)

1.3.11 Clean and Wrangle

Steps are exactly the same as previous sections as dataframe layout identical

In [41]: *#Remove all years from both life_exp and income so they match.*

#Removing years for life_exp

```

col = np.r_[2021:2101]
for i in col:
    col_drop = str(i)
    life_exp_df.drop(col_drop, axis=1, inplace=True)

#Removing years for income
col = np.r_[2021:2041]
for i in col:
    col_drop = str(i)
    income_df.drop(col_drop, axis=1, inplace=True)

#Melt the columns into rows
income_melt_df = pd.melt(income_df, id_vars='country', var_name='year', value_name='income')
life_exp_melt_df = pd.melt(life_exp_df, id_vars='country', var_name='year', value_name='life_exp')

In [42]: #Compare shape of the dataframes to make sure there are 3 columns
income_melt_df.shape

Out[42]: (42653, 3)

In [43]: life_exp_melt_df.shape

Out[43]: (41327, 3)

In [44]: #income has more countries so merge income to life_exp only keeping similar countries
life_income_df = pd.merge(income_melt_df, life_exp_melt_df, on=['country', 'year'], how='inner')
life_income_df.shape

#In truth you could merge into smaller dataframe or do an inner merge here.

Out[44]: (41327, 4)

In [45]: #Check for null values
life_income_df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 41327 entries, 0 to 41326
Data columns (total 4 columns):
country          41327 non-null object
year             41327 non-null object
income_per_capita 41327 non-null int64
exp_years_lived   40808 non-null float64
dtypes: float64(1), int64(1), object(2)
memory usage: 1.6+ MB

```

After checking for null or missing data we find that `exp_years_lived` has lots of missing values. Let's explore more to better understand where the missing values are coming from.


```
In [46]: #Figure out where the missing data is in exp_years_lived.
missing_df = life_income_df.query('exp_years_lived == "NaN"')
missing_df.groupby('country', as_index=False).mean()
```

```
Out[46]:
```

	country	income_per_capita	exp_years_lived
0	Andorra	5240.80925	nan
1	Dominica	1272.83237	nan
2	Marshall Islands	816.12139	nan

```
In [47]: #Seems like its only in 3 different countries for practically all of their data points.
a = ['Andorra', 'Dominica', 'Marshall Islands']
for i in a:
    index_num = life_income_df[life_income_df['country'] == i].index
    life_income_df.drop(index_num, inplace=True)

#Check again for missing data
life_income_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 40664 entries, 0 to 41326
Data columns (total 4 columns):
country          40664 non-null object
year             40664 non-null object
income_per_capita 40664 non-null int64
exp_years_lived   40664 non-null float64
dtypes: float64(1), int64(1), object(2)
memory usage: 1.6+ MB
```

```
In [48]: #Double check that countries have been dropped - dataframe should be empty
life_income_df.query('country == "Andorra" | country == "Dominica" | country == "Marsha
```

```
Out[48]: Empty DataFrame
Columns: [country, year, income_per_capita, exp_years_lived]
Index: []
```

1.3.12 Analysis

In order to answer the questions properly, I need to break the years into smaller sections and combine into a separate dataframe. I will do this by splitting into roughly 50 year segments and a more recent 20 year segment so I can find means in those time periods and see which countries have improved the most without worry to much about skewed data.

```
In [49]: #Used this for comparison when printing out the early - late dataframes
life_income_df.groupby('country', as_index=False).mean().head()
```

```
Out[49]:
```

	country	income_per_capita	exp_years_lived
0	Afghanistan	1319.09955	35.30891
1	Albania	2386.40724	47.20498

2	Algeria	4254.33032	41.22534
3	Angola	2254.69231	35.07421
4	Antigua and Barbuda	5102.78281	47.06787

```
In [50]: #First, I want to understand how life expectancy has changed overtime
x = 'exp_years_lived'
```

```
#Removing timeframes from the life_income_df so we can combine into another df later
early_df = life_income_df.query('year < "1850"').groupby('country', as_index=False)[x].
mid_df = life_income_df.query('year >= "1850" & year < "1900"').groupby('country', as_i
midlate_df = life_income_df.query('year >= "1900" & year < "1950"').groupby('country',
late_df = life_income_df.query('year >= "1950" & year < "2000"').groupby('country', as_
current_df = life_income_df.query('year >= "2000"').groupby('country', as_index=False)[

#Check to make sure it worked
current_df.head()
```

```
Out[50]:
```

	country	exp_years_lived
0	Afghanistan	59.57143
1	Albania	77.08571
2	Algeria	76.22857
3	Angola	59.35714
4	Antigua and Barbuda	76.34762

```
In [51]: #Change column names so when combine we can tell them apart
early_df.rename(columns={x : 'early_yrs'}, inplace=True)
mid_df.rename(columns={x : 'mid_yrs'}, inplace=True)
midlate_df.rename(columns={x : 'midlate_yrs'}, inplace=True)
late_df.rename(columns={x : 'late_yrs'}, inplace=True)
current_df.rename(columns={x : 'current_yrs'}, inplace=True)

#Check changes
mid_df.head()
```

```
Out[51]:
```

	country	mid_yrs
0	Afghanistan	28.09200
1	Albania	35.40000
2	Algeria	27.86400
3	Angola	27.58800
4	Antigua and Barbuda	33.58200

```
In [52]: #Combine the time sections for average years lived for all countries
from functools import reduce
data_frames = [early_df, mid_df, midlate_df, late_df, current_df]
merge_df = reduce(lambda left,right: pd.merge(left,right,on=['country'],
                                                how='outer'), data_frames)

#Check merge
merge_df.head()
```

```

Out[52]:
      country  early_yrs  mid_yrs  midlate_yrs  late_yrs  current_yrs
0  Afghanistan    27.96000  28.09200     30.47340  44.52000     59.57143
1    Albania      35.40000  35.40000     37.92800  67.54200     77.08571
2    Algeria      28.62400  27.86400     32.69600  61.01600     76.22857
3    Angola       27.00000  27.58800     30.59200  44.91800     59.35714
4  Antigua and Barbuda  33.50000  33.58200     40.52600  68.36600     76.34762

In [53]: #Add an improvement column by subtracting early from late and sort appropriately.
merge_df['improvement'] = merge_df['current_yrs'] - merge_df['early_yrs']
merge_df.sort_values(by='improvement', ascending=False).head(10)

```

```

Out[53]:
      country  early_yrs  mid_yrs  midlate_yrs  late_yrs  current_yrs  \
86    Kuwait    26.00000  26.10800     29.19400  68.86000     81.11905
149  South Korea  25.80000  25.83800     33.19800  62.68600     80.32857
151    Spain     29.50000  30.32600     46.00600  72.63600     81.63810
79    Italy     29.70000  33.53000     49.98800  72.78600     81.79524
116  Nicaragua  25.40000  25.54000     30.17000  59.62600     77.10476
101    Malta     28.70000  32.26000     52.88000  71.52400     80.21429
143    Singapore  31.84400  34.04600     39.47400  70.45800     82.74286
75    Iran       25.60000  25.80200     28.10800  59.58800     75.43810
168    Tunisia   27.06600  28.47000     32.32200  60.18000     76.80952
78    Israel     32.00000  32.00000     36.37800  72.49800     81.29524

      improvement
86      55.11905
149     54.52857
151     52.13810
79      52.09524
116     51.70476
101     51.51429
143     50.89886
75      49.83810
168     49.74352
78      49.29524

```

Kuwait, South Korea, Spain, Italy and Nicaragua have the greatest improvement in life expectancy in the past 220 years based on the data given by Gapminder. Again, I used time segments to steer clear of recently or earlier skewed data. I believe this is a more appropriate representation of improvement in life expectation.

1.3.13 Analysis Cont.

Next, we need to tackle how income per person effects life expectancy. As discussed above, the figures have already been converted/are represented in USD. And again, income in average per person in the given country.

```

In [54]: life_income_df.head()

Out[54]:
      country  year  income_per_capita  exp_years_lived
0  Afghanistan  1800                603         28.20000

```

1	Albania	1800	667	35.40000
2	Algeria	1800	715	28.80000
4	Angola	1800	618	27.00000
5	Antigua and Barbuda	1800	757	33.50000

Let's create the same time segments again and compare income per person to life expectancy in all timeline to see how and if the correlation has changed overtime.

In [55]: *#Create the timed segments so we can merge into one table.*

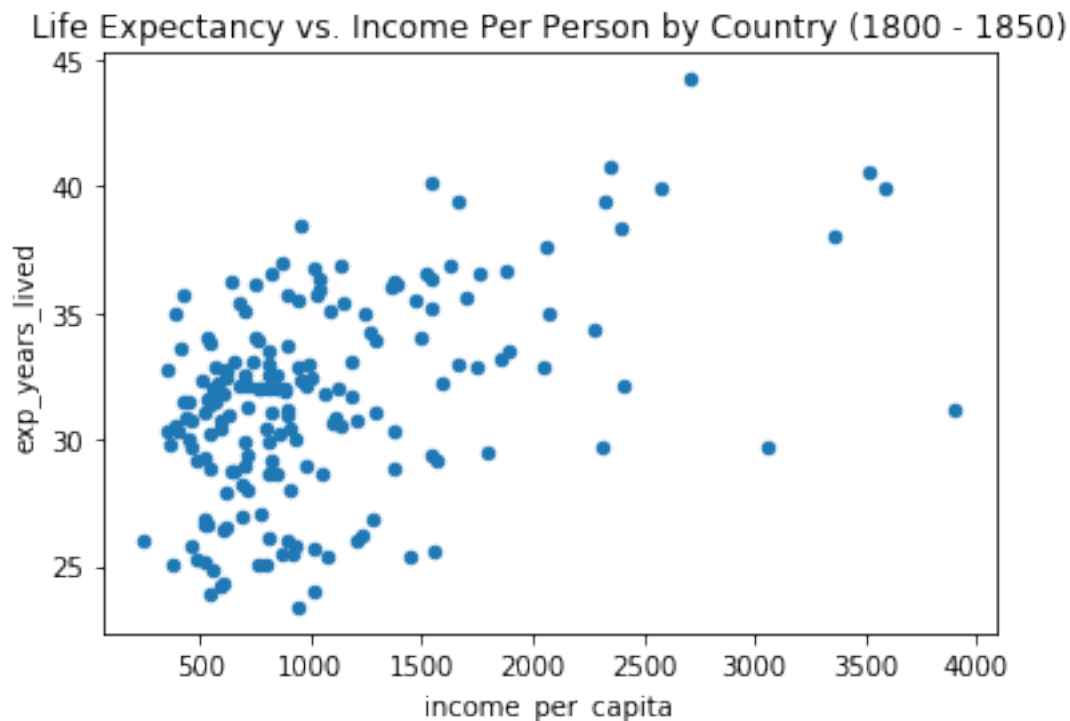
```
early_df = life_income_df.query('year < "1850"').groupby('country', as_index=False).mean()
mid_df = life_income_df.query('year >= "1850" & year < "1900"').groupby('country', as_index=False).mean()
midlate_df = life_income_df.query('year >= "1900" & year < "1950"').groupby('country', as_index=False).mean()
late_df = life_income_df.query('year >= "1950"').groupby('country', as_index=False).mean()
current_df = life_income_df.query('year >= "2000"').groupby('country', as_index=False).mean()
```

#Create some scatter plots so we can compare different timeframes


```
print('Below are scatter plots of expected years to live compared to income per person')
early_df.plot(x="income_per_capita", y="exp_years_lived", kind='scatter', title='Life Expectancy vs. Income Per Person by Country (1800 - 1850)')
mid_df.plot(x="income_per_capita", y="exp_years_lived", kind='scatter', title='Life Expectancy vs. Income Per Person by Country (1850 - 1900)')
midlate_df.plot(x="income_per_capita", y="exp_years_lived", kind='scatter', title='Life Expectancy vs. Income Per Person by Country (1900 - 1950)')
late_df.plot(x="income_per_capita", y="exp_years_lived", kind='scatter', title='Life Expectancy vs. Income Per Person by Country (1950 - 2000)')
current_df.plot(x="income_per_capita", y="exp_years_lived", kind='scatter', title='Life Expectancy vs. Income Per Person by Country (2000 - Present)')
```

Below are scatter plots of expected years to live compared to income per person organized by country


Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x7f64cf986eb8>




Life Expectancy vs. Income Per Person by Country (1850 - 1900)




Life Expectancy vs. Income Per Person by Country (1900 - 1950)



Life Expectancy vs. Income Per Person by Country (1950 - 2000)



Life Expectancy vs. Income Per Person by Country (2000 - 2020)



After scanning the scatter plots it almost seems like the correlation between and income per person and longer life expectancy has become less and less correlated. To be sure lets use the Pearson correlation coefficient again on each timeline.

```
In [56]: #Create an array of data for correlation coefficients
        #NOTE: I had to give dataframes below names so I could append the names appropriately i
        early_df.name = 'early'
        mid_df.name = 'mid'
        midlate_df.name = 'midlate'
        late_df.name = 'late'
        current_df.name = 'current'
        df = [early_df, mid_df, midlate_df, late_df, current_df]
        data = []
        for i in df:
            column_1 = i["income_per_capita"]
            column_2 = i["exp_years_lived"]
            correlation = (column_1.corr(column_2)) * 100
            data.append([i.name, correlation])

        #Use the array of data to create a new data frame to anaylze.
        corr_df = pd.DataFrame(data, columns = ['Era', 'Coefficient'])
        corr_df
```

```
Out[56]:
```

	Era	Coefficient
0	early	45.96669
1	mid	58.32707
2	midlate	75.68726
3	late	54.55680
4	current	61.56485

1.3.14 Q3 Synopsis

Based on our above analysis it seems that the most improved countries on life expectancy are Kuwait, South Korea, Spain, Italy and Nicaragua.

As for the correlation between life expectancy the average income per person, there seems to be little correlation but not enough to make a stronger agrument.

Although it doesn't seem like income has a strong correlation or effect on life expectancy within countries I wonder if there are of dependent variable that have an effect. Maybe health spend?

Research Question 4: How health spend effect life expectancy?

1.3.15 Explanation

I was unlikely in the comparison between income per person and life expectancy within a country but it would be interesting to see what the correlation between health spend and life expectancy looks like. Also, the health spend is the % of total GDP which I think is a good metric because we will be looking at proportions rather than hard numbers per country.

As I'm performing many of the same functions I'm going to follow the same basic order within this question and possibly use dataframes from the past question.

```
In [57]: #Libraries and dataset:
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
% matplotlib inline

health_df = pd.read_csv('total_health_spend.csv')
life_exp_df = pd.read_csv('life_exp.csv')
```

1.3.16 Understanding Data Layout

Based on below functions around the data it is layed out identical to all other Gapminder data. We will need to remove irrelevant years, melt the data and combine reletive dataframes.

```
In [58]: #Unfortunately health spend as % of total GDP only goes from 1995 to 2010
health_df.head()
```

```
Out[58]:
```

	country	1995	1996	1997	1998	1999	2000	2001	\
0	Afghanistan	nan	nan	nan	nan	nan	nan	nan	
1	Albania	0.02560	0.04040	0.04810	0.05340	0.05830	0.06350	0.06040	
2	Algeria	0.04170	0.03790	0.04060	0.04060	0.03870	0.03490	0.03840	
3	Andorra	0.07640	0.08030	0.08030	0.09850	0.07350	0.07580	0.06790	
4	Angola	0.03790	0.02170	0.02260	0.01850	0.01940	0.02410	0.03450	

		2002	2003	2004	2005	2006	2007	2008	2009	2010
0		0.05720	0.06820	0.06360	0.06630	0.06770	0.07300	0.06980	0.07580	0.07580
1		0.06280	0.06160	0.06880	0.06840	0.06730	0.06880	0.06750	0.06880	0.06550
2		0.03870	0.03740	0.03380	0.03060	0.03130	0.03530	0.03730	0.04580	0.04170
3		0.07040	0.07120	0.07110	0.07220	0.07440	0.07550	0.07540	0.07520	0.07520
4		0.02360	0.02640	0.02090	0.01960	0.02420	0.02520	0.03190	0.04950	0.02850

```
In [59]: #remove years for life_exp to match health spend (1996 - 2010)
col = np.r_[1800:1996, 2011:2101]
for i in col:
    col_drop = str(i)
    life_exp_df.drop(col_drop, axis=1, inplace=True)

#melt the columns into rows
health_melt_df = pd.melt(health_df, id_vars='country', var_name='year', value_name='health_spend')
life_exp_melt_df = pd.melt(life_exp_df, id_vars='country', var_name='year', value_name='life_exp')
```

```
In [60]: #Make sure code ran successfully
life_exp_melt_df.head()
```

```
Out[60]:
```

	country	year	exp_years_lived
0	Afghanistan	1996	53.80000

1	Albania	1996	74.40000
2	Algeria	1996	73.30000
3	Andorra	1996	80.00000
4	Angola	1996	50.10000

```
In [61]: #Double check to make sure there are 3 rows and understand more fully how to merge the
health_melt_df.shape
```

```
Out[61]: (3040, 3)
```

```
In [62]: life_exp_melt_df.shape
```

```
Out[62]: (2805, 3)
```

```
In [63]: life_health_df = pd.merge(health_melt_df, life_exp_melt_df, on=['country', 'year'], how
life_health_df.shape
```

```
Out[63]: (2760, 4)
```

```
In [64]: #inspect the data to further check
life_health_df.head()
```

```
Out[64]:
```

	country	year	health_spend	exp_years_lived
0	Afghanistan	1996	nan	53.80000
1	Albania	1996	0.04040	74.40000
2	Algeria	1996	0.03790	73.30000
3	Andorra	1996	0.08030	80.00000
4	Angola	1996	0.02170	50.10000

```
In [65]: #Check for missing data
life_health_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2760 entries, 0 to 2759
Data columns (total 4 columns):
country          2760 non-null object
year             2760 non-null object
health_spend     2731 non-null float64
exp_years_lived  2760 non-null float64
dtypes: float64(2), object(2)
memory usage: 107.8+ KB
```

We are missing some health_spend values! Next, we will have to see where the data is missing from and determine how to fix

```
In [66]: #understand data so I can clean
life_health_df.query('health_spend == "NaN"')
```

```
Out[66]:
```

	country	year	health_spend	exp_years_lived
0	Afghanistan	1996	nan	53.80000
94	Liberia	1996	nan	48.90000
184	Afghanistan	1997	nan	53.70000
278	Liberia	1997	nan	51.90000
368	Afghanistan	1998	nan	52.80000
552	Afghanistan	1999	nan	54.40000
736	Afghanistan	2000	nan	54.60000
920	Afghanistan	2001	nan	54.80000
1252	Somalia	2002	nan	53.30000
1287	Zimbabwe	2002	nan	45.60000
1436	Somalia	2003	nan	53.70000
1471	Zimbabwe	2003	nan	45.40000
1620	Somalia	2004	nan	54.00000
1655	Zimbabwe	2004	nan	45.10000
1804	Somalia	2005	nan	54.70000
1839	Zimbabwe	2005	nan	45.10000
1988	Somalia	2006	nan	55.10000
2023	Zimbabwe	2006	nan	45.40000
2172	Somalia	2007	nan	55.00000
2207	Zimbabwe	2007	nan	45.90000
2356	Somalia	2008	nan	55.50000
2391	Zimbabwe	2008	nan	46.30000
2540	Somalia	2009	nan	55.90000
2575	Zimbabwe	2009	nan	47.20000
2648	Honduras	2010	nan	72.80000
2683	Mexico	2010	nan	75.20000
2695	Nicaragua	2010	nan	77.60000
2724	Somalia	2010	nan	55.00000
2759	Zimbabwe	2010	nan	49.70000

```
In [67]: #Zimbabwe, Somalia and Afghanistan have to many data points missing so drop them
a = ['Zimbabwe', 'Somalia', 'Afghanistan']
for i in a:
    index_nums = life_health_df[life_health_df['country'] == i].index
    life_health_df.drop(index_nums, inplace=True)

#Check to see if worked
life_health_df.query('health_spend == "NaN"')
```

```
Out[67]:
```

	country	year	health_spend	exp_years_lived
94	Liberia	1996	nan	48.90000
278	Liberia	1997	nan	51.90000
2648	Honduras	2010	nan	72.80000
2683	Mexico	2010	nan	75.20000
2695	Nicaragua	2010	nan	77.60000

```
In [68]: #For all other countries I'll find mean and replace the corresponding row values with s
```

```

#Find mean for all the countries health spend
liberia_mean = life_health_df.query('country == "Liberia"')['health_spend'].mean()
honduras_mean = life_health_df.query('country == "Honduras"')['health_spend'].mean()
mexico_mean = life_health_df.query('country == "Mexico"')['health_spend'].mean()
nicaragua_mean = life_health_df.query('country == "Nicaragua"')['health_spend'].mean()

#replace the NaN inputs with the mean values with the specific rows
life_health_df['health_spend'] = np.where((life_health_df['country'] == 'Liberia') & (life_health_df['health_spend'].isnull()), liberia_mean, life_health_df['health_spend'])
life_health_df['health_spend'] = np.where((life_health_df['country'] == 'Honduras') & (life_health_df['health_spend'].isnull()), honduras_mean, life_health_df['health_spend'])
life_health_df['health_spend'] = np.where((life_health_df['country'] == 'Mexico') & (life_health_df['health_spend'].isnull()), mexico_mean, life_health_df['health_spend'])
life_health_df['health_spend'] = np.where((life_health_df['country'] == 'Nicaragua') & (life_health_df['health_spend'].isnull()), nicaragua_mean, life_health_df['health_spend'])

```

In [69]: #check to make sure data is cleaned

```

life_health_df.info()
life_health_df.query('health_spend == 0')

#Make sure Honduras has every year listed (1996 - 2010)
life_health_df.query('country == "Honduras"')

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2715 entries, 1 to 2758
Data columns (total 4 columns):
country      2715 non-null object
year         2715 non-null object
health_spend 2715 non-null float64
exp_years_lived 2715 non-null float64
dtypes: float64(2), object(2)
memory usage: 106.1+ KB

```

```

Out[69]:
   country  year  health_spend  exp_years_lived
72  Honduras  1996      0.04960      69.40000
256 Honduras  1997      0.04880      69.70000
440 Honduras  1998      0.05630      63.50000
624 Honduras  1999      0.05700      70.10000
808 Honduras  2000      0.05380      70.40000
992 Honduras  2001      0.05450      70.70000
1176 Honduras 2002      0.06070      71.00000
1360 Honduras 2003      0.06400      71.20000
1544 Honduras 2004      0.06470      71.40000
1728 Honduras 2005      0.05830      71.70000
1912 Honduras 2006      0.05720      71.90000
2096 Honduras 2007      0.05910      72.20000
2280 Honduras 2008      0.06050      72.50000
2464 Honduras 2009      0.07000      72.70000
2648 Honduras 2010      0.05818      72.80000

```


1.3.17 Analysis

For here it is very easy compared to some of our other questions. Since we are only working with a small segment of time we can go to take the average across entire time period and compare the variables in a scatter plot to better spot a correlation.

Furthermore, I will find the correlation coefficient to compare our earlier correlation of income per person to life expectancy to our new correlation using health spend.

```
In [70]: #Create scatter plot using mean values for countries
life_health_df.groupby('country', as_index=False).mean().plot(x='health_spend', y='exp_years_lived', title='Life Expectancy vs Health Spend by Country')
```

```
Out[70]: <matplotlib.axes._subplots.AxesSubplot at 0x7f64cf86d9b0>
```



```
In [71]: #Pearson Correlation Coefficient - higher the coefficient the more correlated the data
column_1 = life_health_df["health_spend"]
column_2 = life_health_df["exp_years_lived"]
correlation = column_1.corr(column_2)
print(correlation * 100)
```

```
24.5121296231
```

1.3.18 Q4 Synapsis

It is fair to say that there is no correlation between health spend (% of GDP) and life expectancy by country. There is only about a 24.5% correlation between the two and we can't argue for a correlation until at least after 70%.

Project Final Conclusion

Results: 1. When comparing violent crimes (total murders) in proportion to countries population, Colombia, El Salvador, Guatemala, South Africa and Venezuela are the 5 countries with the highest statistical probability of violent attacks per person. 2. There is a -2.55% correlation between population density per country and total murders. Based on the data parameters, there is no correlation and population density doesn't effect total murders by country. 3. The top 5 most improved countries from early 1800's to post-2000 for average years lived per person are Kuwait, South Korea, Spain, Italy and Nicaragua. 4. The correlation between income per capita and life expectancy by country ranges from 45% to 75% from 1800 - 2020. As the correlations are intermittent throughout history and not substantial enough for a strong argument it's safe to say there is little to no correlation. 5. There is only a 24.5% correlation between health spend as a percentage of GDP and average life expectancy in years by country. Based on statistics, there is no correlation between budgetted health spend and life expectancy.

Data Limitations: 1. We are limited to the data provided by Gapminder which does not include all countires and has no murder/population data on individual metropolitan areas which would have helped narrow down most dangerous 'place' to live. 2. Density was measured for entire area (sq km) of country which would skew that data for countries with more uninhabited land. This would be a more useful finding if it was narrowed down to cities. 3. For both total murders and health spend data files, we were limited to under 30 years of availabe data points and more limited on listed countries. Although this may be substantial enough to make a claim it would be as strong as other dataframes and results provided. 4. Global data provided by many diffing resources usually results in differing data integrity and can always be challenged and questioned.

```
In [ ]: from subprocess import call
        call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```